

Introduction to Apache Kafka for Developers



Apache Kafka is an open-source distributed event streaming platform used by thousands of companies for high-performance data pipelines, streaming analytics, data integration, and mission-critical applications.

Lab files

Every lab folder has a `README.md` with instructions on how to complete the lab, and a `solution` folder with a working version of the lab for reference. Please use the README.md and solution folder as a guide to completing the labs.

Lab Instructions

The latest instructions for the labs can be found below.

1. [Lab 1](#) Basic Kafka
2. [Lab 1.2](#) Understanding Failover
3. [Lab 2](#) Writing a simple Java Producer
4. [Lab 3](#) Writing a simple Java Consumer
5. [Lab 5.1](#) Advanced Producers
6. [Lab 5.2](#) Clean shutdown
7. [Lab 5.3](#) Configuring Producer Durability
8. [Lab 5.4](#) Metrics and Replication Verification
9. [Lab 5.5](#)
10. [Lab 5.6](#) Retry, timeouts, in-flight, back-off
11. [Lab 5.7](#)
12. [Lab 5.8](#)
13. [Lab 6.1](#) Advanced Consumer
14. [Lab 6.2](#) Position Consumer in Topic Log
15. [Lab 6.3](#) At-Most-Once and At-Least-Once Semantics
16. [Lab 6.4](#) Exactly-Once message semantics
17. [Lab 6.5](#) Thread per consumer
18. [Lab 6.6](#) Consumer with worker threads
19. [Lab 6.7](#) Priority queue with consumer.assign
20. [Lab Streams 1](#) Kafka Streams Lab 1
21. [Lab Streams 2](#) Kafka Streams Lab 2