

LAB Stream 1-2 Custom Word Count

Example

Welcome to the session Stream 1-2 lab. The work for this lab is done in `~/kafka-training/labs/stream-lab1-2`. In this lab, you are going to run a custom Word Count Example. You must write the java code You create two new Kafka topics, one called `word-count-input` and the other `word-count-output`. You will then add some data to the input topic, start a consumer on the output topic so you can watch it, and then run the Word Count Demo

Note: Lab solution is available in following directory:

```
~/kafka-training/labs/stream-lab1-2/solution
```

Create Kafka Topics

- Create a topic named `word-count-input` with 1 partition and a replication factor of 1.
- Create a topic named `word-count-output` with 1 partition and a replication factor of 1.

ACTION - Edit `create-topics.sh`

`~/kafka-training/labs/stream-lab1-2/bin/create-topics.sh`

```
#!/usr/bin/env bash
cd ~/kafka-training

## Create input topic
kafka/bin/kafka-topics.sh --create \
  --replication-factor 1 \
  --partitions 1 \
  --topic word-count-input \
  --zookeeper localhost:2181

## Create output topic
kafka/bin/kafka-topics.sh --create \
  --replication-factor 1 \
  --partitions 1 \
  --topic word-count-output \
  --zookeeper localhost:2181

## List created topics
kafka/bin/kafka-topics.sh --list \
  --zookeeper localhost:2181
```

ACTION - START ZooKeeper and Kafka Broker if needed. Only 1 broker is necessary.

ACTION - RUN `create-topics.sh` as follows:

```
$ cd ~/kafka-training/labs/stream-lab1-2/bin
$ ./create-topics.sh

Created topic "word-count-input".
```

```
Created topic "word-count-output".
word-count-input
word-count-output
```

Add data to the topic

Now add data to the input topic by starting a console producer and entering data.

ACTION - REVIEW `start-producer-console-input.sh`

`~/kafka-training/labs/stream-lab1-2/bin/start-producer-console-input.sh`

```
#!/usr/bin/env bash
cd ~/kafka-training

## Producer
kafka/bin/kafka-console-producer.sh --broker-list localhost:9092 --topic streams-plaintext-input
```

ACTION - RUN `start-producer-console-input.sh`

```
$ cd ~/kafka-training/labs/stream-lab1-2
$ ./start-producer-console-input.sh
>
```

ACTION - Enter Data

Enter the following 3 lines of data. Ctrl-C to stop the console.

```
> stock streams MSFT
> stock data AAPL
> stock streams RHT
```

ACTION - Validate Data

Run a console consumer against the input topic to verify the data is there.

`~/kafka-training/labs/stream-lab1-2/start-consumer-console-input.sh`

```
#!/usr/bin/env bash
cd ~/kafka-training

## Input Consumer
kafka/bin/kafka-console-consumer.sh --bootstrap-server localhost:9092 --topic streams-plaintext-input --from-beginning
```

ACTION - RUN `start-consumer-console-input.sh`

```
$ cd ~/kafka-training/labs/stream-lab1-2/bin
$ ./start-consumer-console-input.sh
> stock streams MSFT
> stock data AAPL
```

```
> stock streams RHT
<Ctrl-C>
Processed a total of 3 messages
```

Above, stops the consumer.

Run a console consumer against the output.

Run another consumer against the output topic so we can see the work that the Word Count Demo does. Notice a few things about this consumer. These match the way the topic is populated by the demo.

- It has a message formatter specified.
- It has key and value deserializers.

ACTION - EDIT `~/kafka-training/labs/stream-lab1-2/bin/start-consumer-console-output.sh`, follow instructions in file.

~/kafka-training/labs/stream-lab1-2/start-consumer-console-output.sh

```
#!/usr/bin/env bash
cd ~/kafka-training

## Output Consumer
kafka/bin/kafka-console-consumer.sh --bootstrap-server localhost:9092 \
  --topic word-count-output \
  --from-beginning \
  --formatter kafka.tools.DefaultMessageFormatter \
  --property print.key=true \
  --property print.value=true \
  --property
key.deserializer=org.apache.kafka.common.serialization.StringDeserializer \
  --property
value.deserializer=org.apache.kafka.common.serialization.LongDeserializer
```

ACTION - RUN `start-consumer-console-output.sh`

```
$ cd ~/kafka-training/labs/stream-lab1-2/bin
$ ./start-consumer-console-output.sh
```

There will be no output until the demo code runs. Keep this window open so you can see it as the demo runs.

Create and Run the Custom Word Count

ACTION - EDIT `~/kafka-training/labs/stream-lab1-2/WordCount.java`, follow instructions in file.

ACTION - RUN `WordCount` Run the java main.

Consumer Console Output after running run-word-count-demo.sh

```
stock    1
streams  1
msft     1
stock    2
```

```
data      1
aapl      1
stock     3
streams   2
rht       1
```

Conclusion Custom Word Count example

The word count is the simplest example of stream processing of an input topic and writing results to a different then you created Kafka Producer in Java that uses the Kafka replicated topic to send records. You sent records with the Kafka Producer using async and sync send methods.

Review Custom Word Count example

How could you change the code to not be case-insensitive?

Remove the `.toLowerCase()` from step 2

How would you simplify the code by using the builder syntax?

Chain calls together