

Kafka Streams and KTables

The background of the slide is a solid dark blue. In the lower half, there are several thick, wavy, light blue lines that sweep from the left side towards the right, creating a sense of motion or data flow.

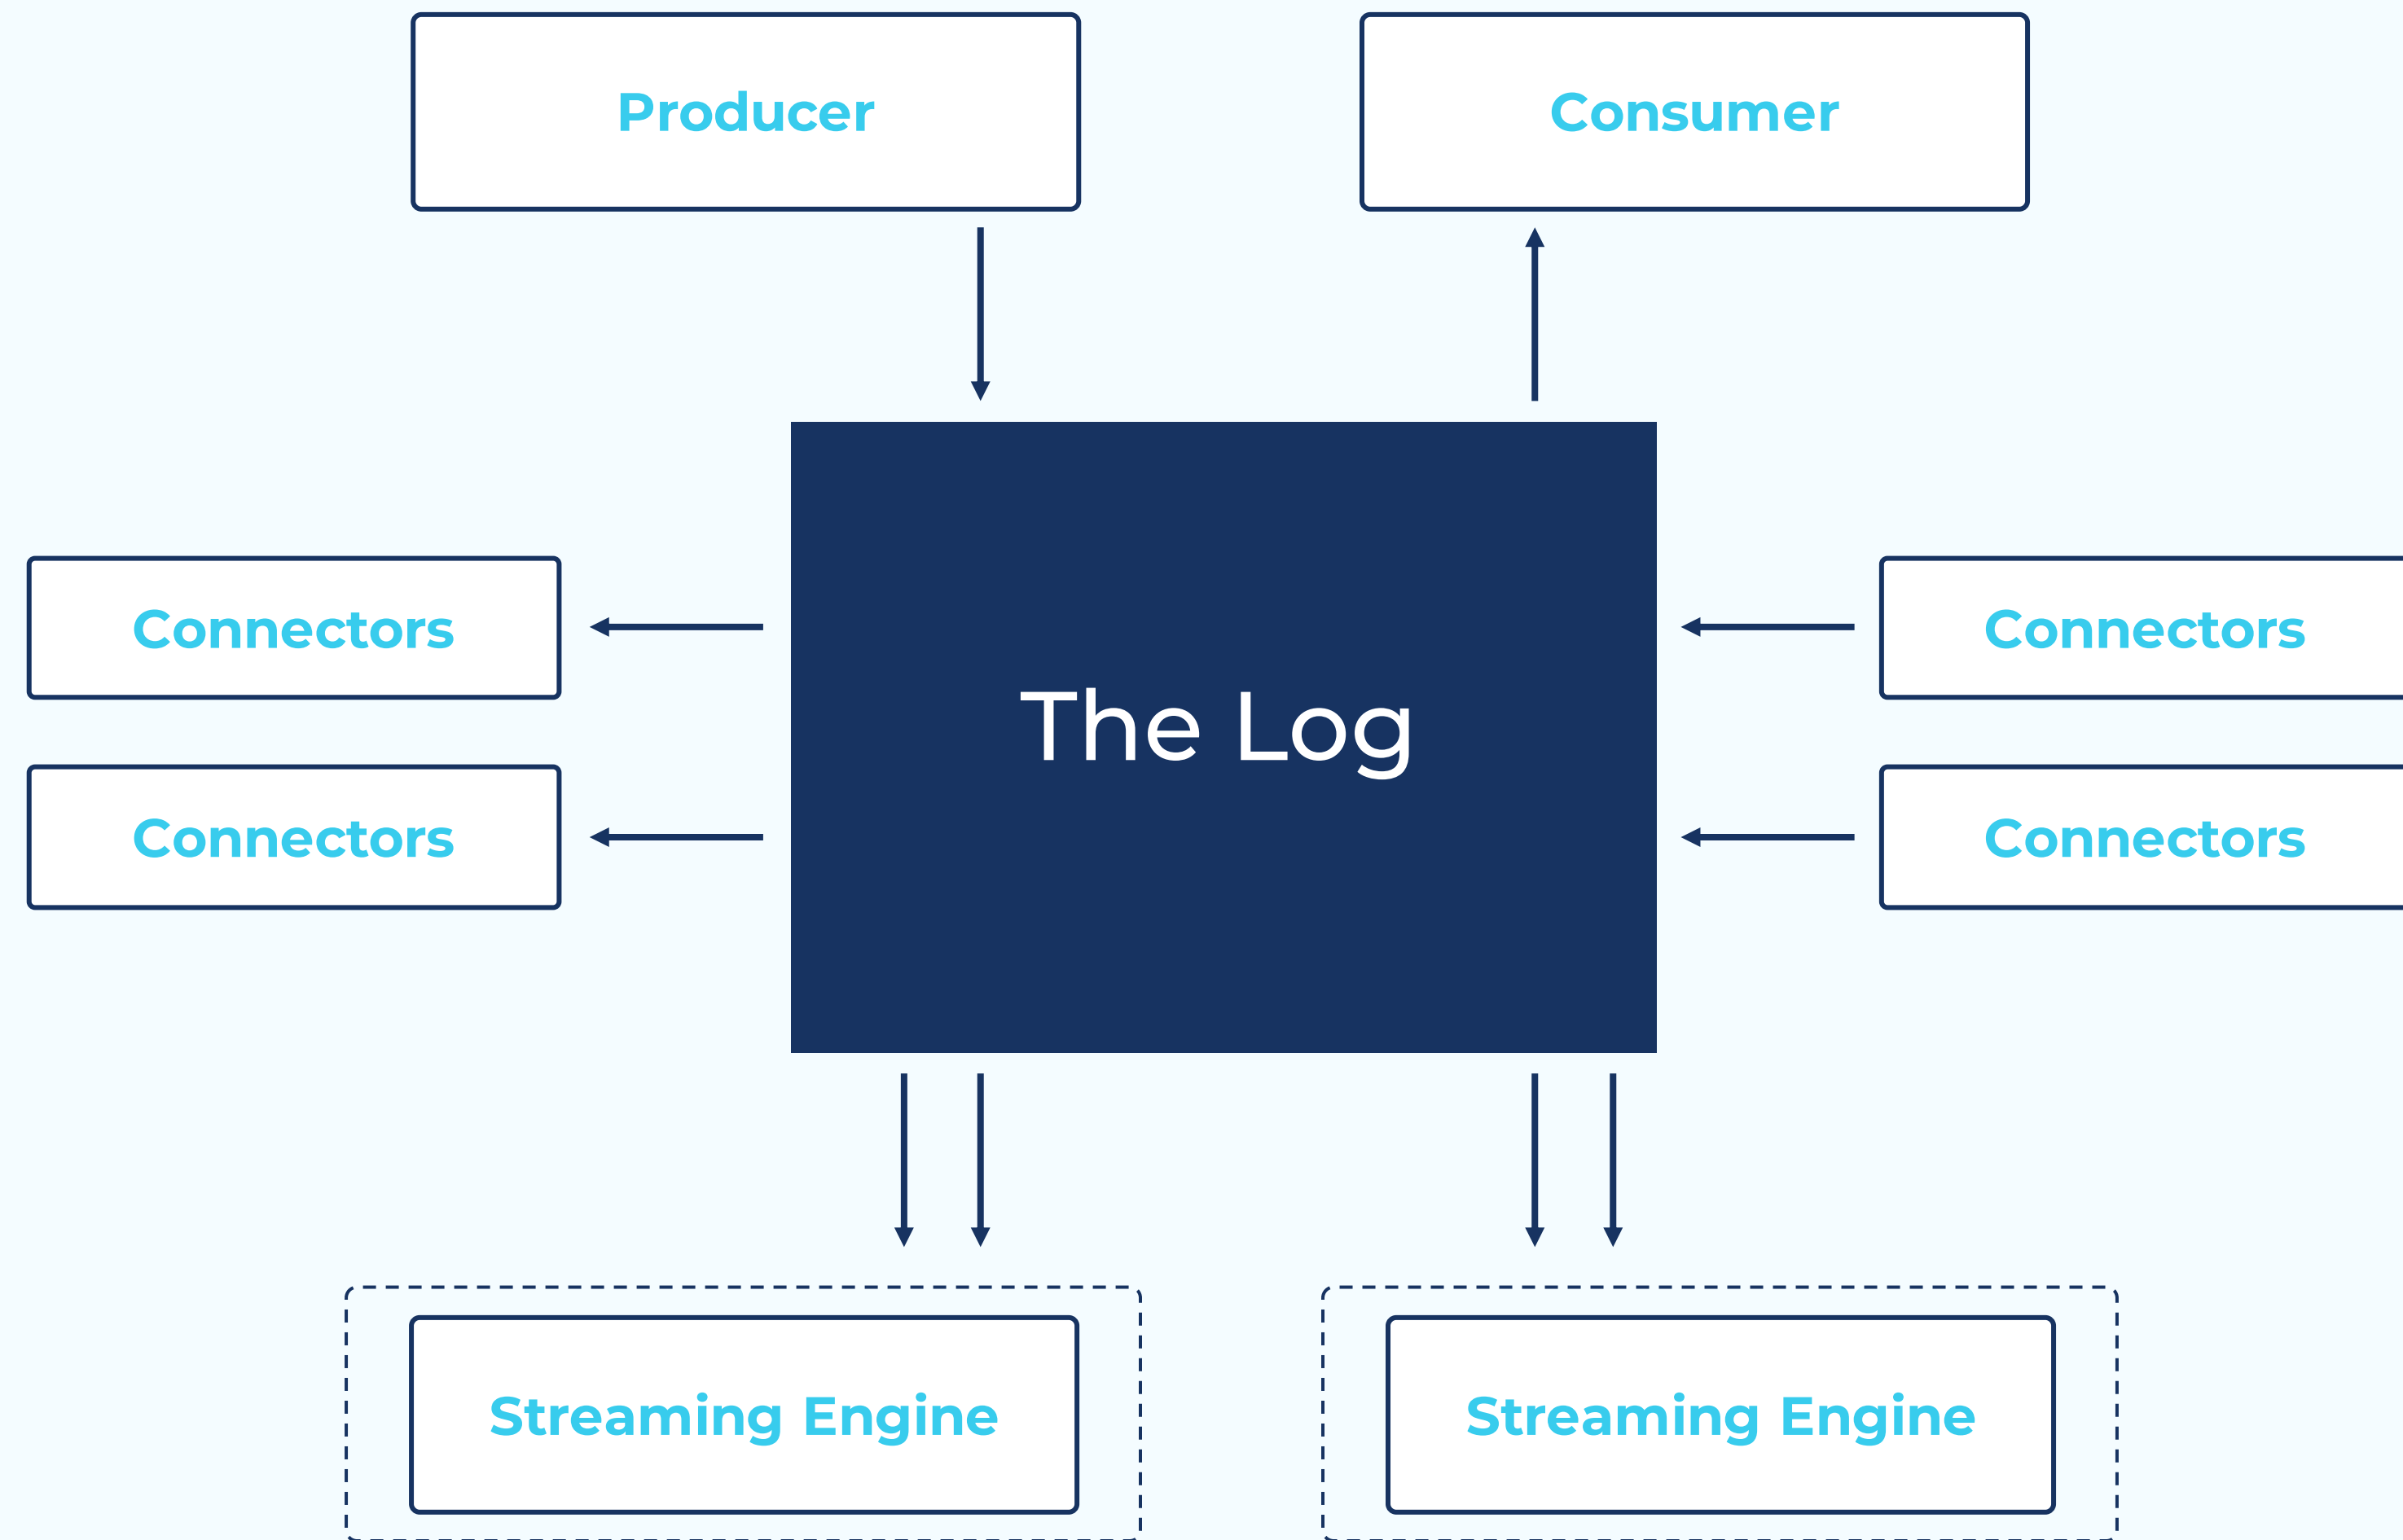
What's Covered

1	Getting Started with Kafka	13	Time Concepts
Streams		14	Hands On: Time Concepts
2	Basic Operations	15	Processor API
3	Hands On: Basic Operations	16	Hands On: Processor API
4	KTable	17	Testing
5	Hands On: KTable	18	Hands On: Testing
6	Serialization	19	Error Handling
7	Joins	20	Hands On: Error Handling
8	Hands On: Joins	21	Internals
9	Stateful Operations	22	Stateful Fault Tolerance
10	Hands On: Aggregations	23	Interactive Queries
11	Windowing		
12	Hands On: Windowing		

Getting Started with Kafka Streams


The bottom half of the slide features several thick, wavy, light blue lines that sweep from the left towards the right, creating a sense of motion and flow against the dark blue background.

Apache Kafka[®]



How to Process Events in Kafka

```
{  
  "reading_ts": "2020-02-14T12:19:27Z",  
  "sensor_id": "aa-101",  
  "production_line": "w01",  
  "widget_type": "acme94",  
  "temp_celcius": 23,  
  "widget_type": 100  
}
```



Processing Events with Producer and Consumer Clients


```
public static void main(String[] args) {  
    try(Consumer<String, Widget> consumer = new KafkaConsumer<>(consumerProperties());  
        Producer<String, Widget> producer = new KafkaProducer<>(producerProperties())) {  
        consumer.subscribe(Collections.singletonList("widgets"));  
        while (true) {  
            ConsumerRecords<String, Widget> records = consumer.poll(Duration.ofSeconds(5));  
            for (ConsumerRecord<String, Widget> record : records) {  
                Widget widget = record.value();  
                if (widget.getColour().equals("red")) {  
                    ProducerRecord<String, Widget> producerRecord = new ProducerRecord<>(  
                        "widgets-red", record.key(), widget);  
                    producer.send(producerRecord, (metadata, exception)-> {.....} );  
                }  
            }  
            ...  
        }  
    }  
}
```

Processing Events with Kafka Streams

```
final StreamsBuilder builder = new StreamsBuilder();

builder.stream("widgets", Consumed.with(stringSerde, widgetsSerde))
    .filter( (key, widget) -> widget.getColour.equals("red"))
    .to("widgets-red", Produced.with(stringSerde, widgetsSerde));
```

Kafka Streams is a Java Library



App Or Microservice



Create, read, process events

Use the library to write standard Java/JVM applications that process data in Kafka.
Kafka Streams makes your applications elastic, distributed, scalable, and fault tolerant.




Basic Operations

The background of the slide is a solid dark blue. In the lower half, there are several thick, wavy, light blue lines that sweep from the left side towards the right, creating a sense of motion and depth.

Event Streams



Processor Topology



Defining a Stream

```
StreamBuilder builder = new StreamBuilder();  
KStream<String, String> firstStream =  
builder.stream(inputTopic, Consumed.with(Serdes.String(), Serdes.String()));
```


Mapping



```
mapValues(value -> value.substring(5))
```

```
map((key, value) -> ..)
```

Filtering



```
filter((key, value) -> Long.parseLong(value) > 1000)
```

Hands On: Basic Operations


The bottom half of the slide features several thick, wavy, light blue lines that sweep upwards from the left towards the right, creating a sense of motion and depth against the dark blue background.

Hands On: Basic Operations

KTable

The bottom of the slide features several thick, wavy, light blue lines that sweep upwards from the left towards the right, creating a dynamic, abstract background element.

Update Streams



Events with the same key are considered updates to previous records with the same key.

Defining a KTable

```
StreamBuilder builder = new StreamBuilder();  
KTable<String, String> firstKTable =  
builder.table(inputTopic, Materialized.with(Serdes.String(), Serdes.String()));
```


Mapping



```
firstKTable.mapValues(value -> ..)
```

```
firstKTable.map((key,value) -> ..)
```

Filtering



```
firstKTable.filter((key, value) -> ..)
```

Global KTable

```
StreamBuilder builder = new StreamBuilder();  
GlobalKTable<String, String> globalKTable =  
Builder.globalTable(inputTopic, Materialized.with(Serdes.String(), Serdes.String()));
```

Hands On: KTable

The bottom half of the slide features several thick, wavy, light blue lines that sweep upwards from the left towards the right, creating a sense of motion and modern design.


Hands On: KTable

Data Serialization


The background of the slide is a dark navy blue. In the lower half, there are several thick, flowing, wavy lines in a lighter blue color that sweep from the left towards the right, creating a sense of motion and depth.

Data Serialization

Bytes to object -
deserialize



Object to bytes -
serialize



Data Serialization - Serdes (Serializer/Deserializer)

```
StreamsBuilder builder = new StreamsBuilder();  
KStream<String, MyObject> stream =  
builder.stream("topic", Consumed.with(Serdes.String(), customObjectSerde))
```

Serdes used by Kafka Streams for
reading input bytes into expected
object types



Data Serialization - Serdes (Serializer/Deserializer)

```
KStream<String, CustomObject> modifiedStream =  
    stream.filter( (key, value) -> value.startsWith("ID5"))  
    .mapValues( value -> new CustomObject(value));  
  
modifiedStream.to("output-topic", Produced.with(Serdes.String(), customObjectSerde));
```



Serdes used by Kafka Streams
for serializing objects into bytes

Data Serialization - Serdes (Serializer/Deserializer)

Data Serialization - Serdes (Serializer/Deserializer)

- Pre-existing serdes: String, Integer, Double, Long, Float, Bytes, ByteArray, ByteBuffer, UUID, and Void
- Additional Serdes available for working with Avro, Protobuf, and JSONSchema
 - Avro
 - SpecificAvroSerde
 - GenericAvroSerde
 - Protobuf
 - KafkaProtobufSerde
 - JSONSchema
 - KafkaJsonSchemaSerde

Joins

The background of the slide is a dark navy blue. In the lower half, there are several thick, wavy, light blue lines that flow from the left side towards the right, creating a sense of movement and depth. The word "Joins" is written in a white, bold, italicized serif font, positioned on the left side of the slide.

Joins

- Kafka Streams offers join operations
- Stream-Stream joins
 - Combine two event streams into a new event stream
 - Join of events based on a common key
 - Records arrive within a defined window of time
 - Possible to compute a new value type
 - Keys are available in read-only mode can be used in computing the new value
- Stream-Table joins
 - KStream-KTable
 - KStream-GlobalKTable
- Table-Table joins

Joins - Types Available

- Stream-Stream
 - Inner - Only if both sides are available within the defined window is a joined result emitted
 - Outer - Both sides always produce an output record
 - Left-value + Right-value
 - Left-value + Null
 - Null + Right-value
 - Left-Outer - The left side always produces an output record
 - Left-value + Right-value
 - Left-value + Null

Joins - Types Available

- Stream-Table
 - Non-windowed joins
 - Inner - Only if both sides are available is a record emitted
 - Left-Outer - The left side (KStream) always produces an output record
 - Left-value + Right-value
 - Left-value + Null
 - Only the Stream side drives the join - new records arriving to the table (right-side) don't result in outputting a join result
 - Applies to both KTable and GlobalKTable joins
 - GlobalKTable joins provide mechanism for determining the join-key from the Stream side key and/or value
 - KTables are timestamp driven but GlobalKTables are bootstrapped - results in different join semantics

Joins - Example

```
KStream<String, String> leftStream = builder.stream("topic-A");
KStream<String, String> rightStream = builder.stream("topic-B");

ValueJoiner<String, String, String> valueJoiner = (leftValue, rightValue) -> {
    return leftValue + rightValue;
};

leftStream.join(rightStream,
               valueJoiner,
               JoinWindows.of(Duration.ofSeconds(10)));
```

Hands On: Joins

The bottom half of the slide features several thick, wavy, light blue lines that sweep across the dark blue background, creating a sense of motion and depth.

Hands On: Joins

Stateful Operations

The background of the slide features several thick, wavy, light blue lines that sweep across the lower half of the frame from left to right, creating a sense of motion and depth against the dark blue background.

Stateful Operations

- Stateless operations are great for operations where you don't need to remember
 - Filter - drop records that don't match a condition
- Other times you need to remember previous results
 - How many times has a particular customer logged in?
 - What's the total sum of tickets sold?
- For those situations where the previous state of an event is important, Kafka Streams offers stateful operations

Stateful Operations - Reduce

```
StreamsBuilder builder = new StreamsBuilder();  
KStream<String, Long> myStream = builder.stream("topic-A");  
  
Reducer<Long> reducer = (longValueOne, longValueTwo) -> longValueOne + longValueTwo;  
myStream.groupByKey().reduce(reducer,  
                             Materialized.with(Serdes.String(), Serdes.Long()))  
        .toStream().to("output-topic");
```

Stateful Operations - Aggregation

```
StreamsBuilder builder = new StreamsBuilder();  
KStream<String, String> myStream = builder.stream("topic-A");  
  
Aggregator<String, String, Long> characterCountAgg =  
    (key, value, charCount) -> value.length() + charCount;  
myStream.groupByKey().aggregate(() -> 0L,  
                                characterCountAgg,  
                                Materialized.with(Serdes.String(),  
Serdes.Long()))  
    .toStream().to("output-topic");
```

Stateful Operations - Considerations

- In Kafka Streams stateful operations don't emit results immediately
- Internal caching buffer results
 - Factors controlling when cache emits records
 - Cache is full (10MB by default)
 - Commit interval (30 seconds)
 - To see all updates, set cache size to zero (also for debugging)

Hands On: Aggregations

The background of the slide is a solid dark blue. In the lower half, there are several thick, wavy, light blue lines that sweep from the left side towards the right, creating a sense of movement and depth. The lines vary in thickness and curvature, with some starting near the bottom left and others higher up, all generally trending upwards and to the right.

Hands On: Aggregations

Windowing

The background of the slide is a dark navy blue. In the lower half, there are several thick, wavy, light blue lines that sweep from the left towards the right, creating a sense of motion and depth. The word 'Windowing' is positioned in the upper left quadrant, rendered in a white, bold, italicized sans-serif typeface.

Windowing

- Aggregations continue to build up over time
- Windowing gives snapshot of an aggregate within a given timeframe
- Four types of windows we'll discuss
 - Tumbling
 - Hopping
 - Session
 - Sliding

Windowing

```
KStream<String, String> myStream = builder.stream("topic-A");  
myStream.groupByKey().count().toStream().to("output")
```


Windowing - Hopping Window

```
KStream<String, String> myStream = builder.stream("topic-A");  
Duration windowSize = Duration.ofMinutes(5);  
Duration advanceSize = Duration.ofMinutes(1);  
TimeWindows hoppingWindow =  
    TimeWindows.of(windowSize).advanceBy(advanceSize);  
myStream.groupByKey()  
    .windowedBy(hoppingWindow)  
    .count();
```

Windowing - Tumbling Window

```
KStream<String, String> myStream = builder.stream("topic-A");  
Duration windowSize = Duration.ofSeconds(30);  
TimeWindows tumblingWindow = TimeWindows.of(windowSize);  
  
myStream.groupByKey()  
        .windowedBy(tumblingWindow)  
        .count();
```

Windowing - Session Window

```
KStream<String, String> myStream = builder.stream("topic-A");  
Duration inactivityGap = Duration.ofMinutes(5);  
  
myStream.groupByKey()  
        .windowedBy(SessionWindows.with(inactivityGap))  
        .count();
```

Windowing - Sliding Window

```
KStream<String, String> myStream = builder.stream("topic-A");  
Duration timeDifference = Duration.ofSeconds(2);  
Duration gracePeriod = Duration.ofMillis(500);  
myStream.groupByKey()  
    .windowedBy(SlidingWindows.withTimeDifferenceAndGrace(  
        timeDifference, gracePeriod))  
    .count();
```

Windowing - Grace Period

```
TimeWindows.of(Duration.ofSeconds(30)).grace(Duration.ofSeconds(5));
```

Hands On: Windowing

The background of the slide is a dark navy blue. In the lower half, there are several thick, wavy, light blue lines that sweep from the left towards the right, creating a sense of motion and depth. The lines vary in thickness and curvature, with some starting near the bottom left and others higher up, all trending towards the right edge of the frame.

Hands On: Windowing

Time Concepts

The background of the slide is a dark navy blue. In the lower half, there are several thick, wavy, light blue lines that flow from the left side towards the right, creating a sense of movement and depth. The lines vary in thickness and curve, with some starting near the bottom left and others higher up, all trending upwards and to the right.

Time Concepts

- Timestamps are a critical component of Kafka.
- The Kafka message format has a dedicated **timestamp** field.
- **Event-time:** A producer, including the Kafka Streams library, automatically sets this timestamp field if user does not
 - Timestamp is the current wall-clock time of the Producer environment when the event is created
- **Ingestion-time:** can configure the Kafka broker to set this timestamp field when an event is appended to (stored in) the topic
 - Timestamp is the current wall-clock time of the Broker environment

Time Concepts

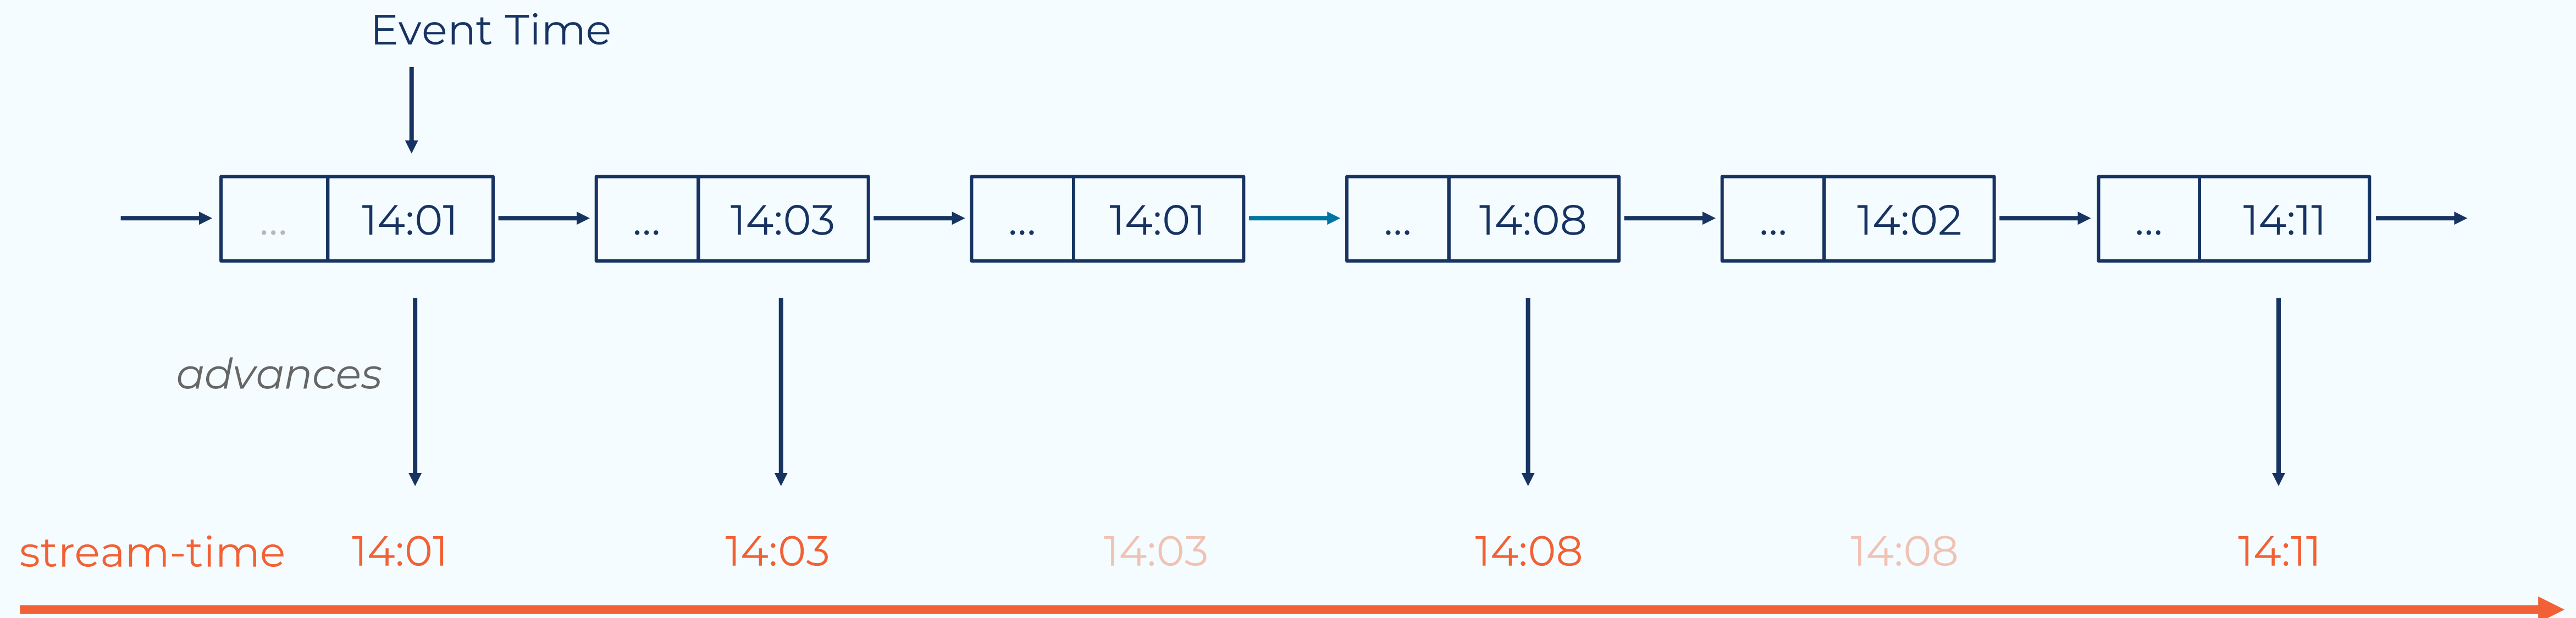
- Timestamps of events drive the action in Kafka Streams
- Earliest timestamp across all input partitions chosen first for processing
- Kafka Streams uses the **TimestampExtractor** interface to get timestamp
 - Default behavior is to use the event timestamp (set by either the event producer or the Kafka broker, see previous slide)
 - Default extractor is the **FailOnInvalidTimestamp**
 - Timestamp set by producer (event-time) or broker (ingestion-time)
- If it's desired to use a timestamp *embedded* in the event “payload” (i.e., the event key or the event value), provide a custom **TimestampExtractor** interface implementation

Time Concepts

- Time moves forward in Kafka Streams by these timestamps.
- For windowing operations this means the timestamps govern the opening and closing of windows.
 - How long a window remains open depends on timestamps only; it's completely detached from wall-clock time
- Kafka Streams has a concept of **Stream Time**

Time Concepts - Stream Time

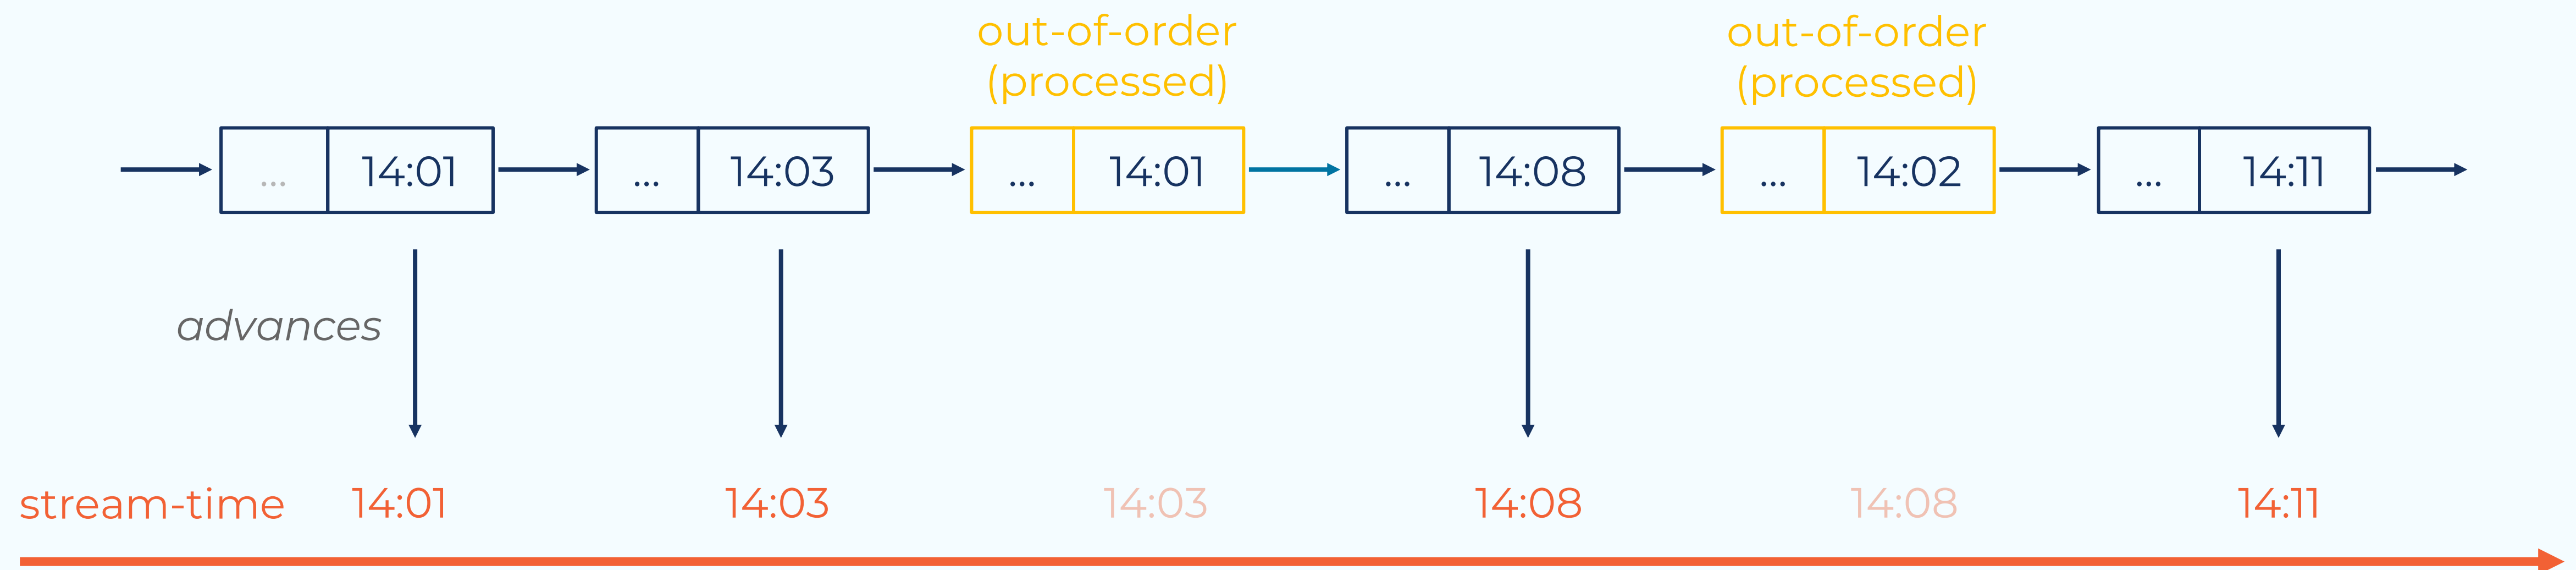
- Stream Time definition (based on “slack time” [1])
 - Largest timestamp seen so far
 - Only moves forward, never backward
 - If an out-of-order event arrives, stream-time stays where it is



[1] cf. [Beyond Analytics: The Evolution of Stream Processing Systems \(SIGMOD 2020\)](#), [Aurora: a new model and architecture for data stream management \(VLDB Journal 2003\)](#)

Time Concepts - Out-of-Order Input

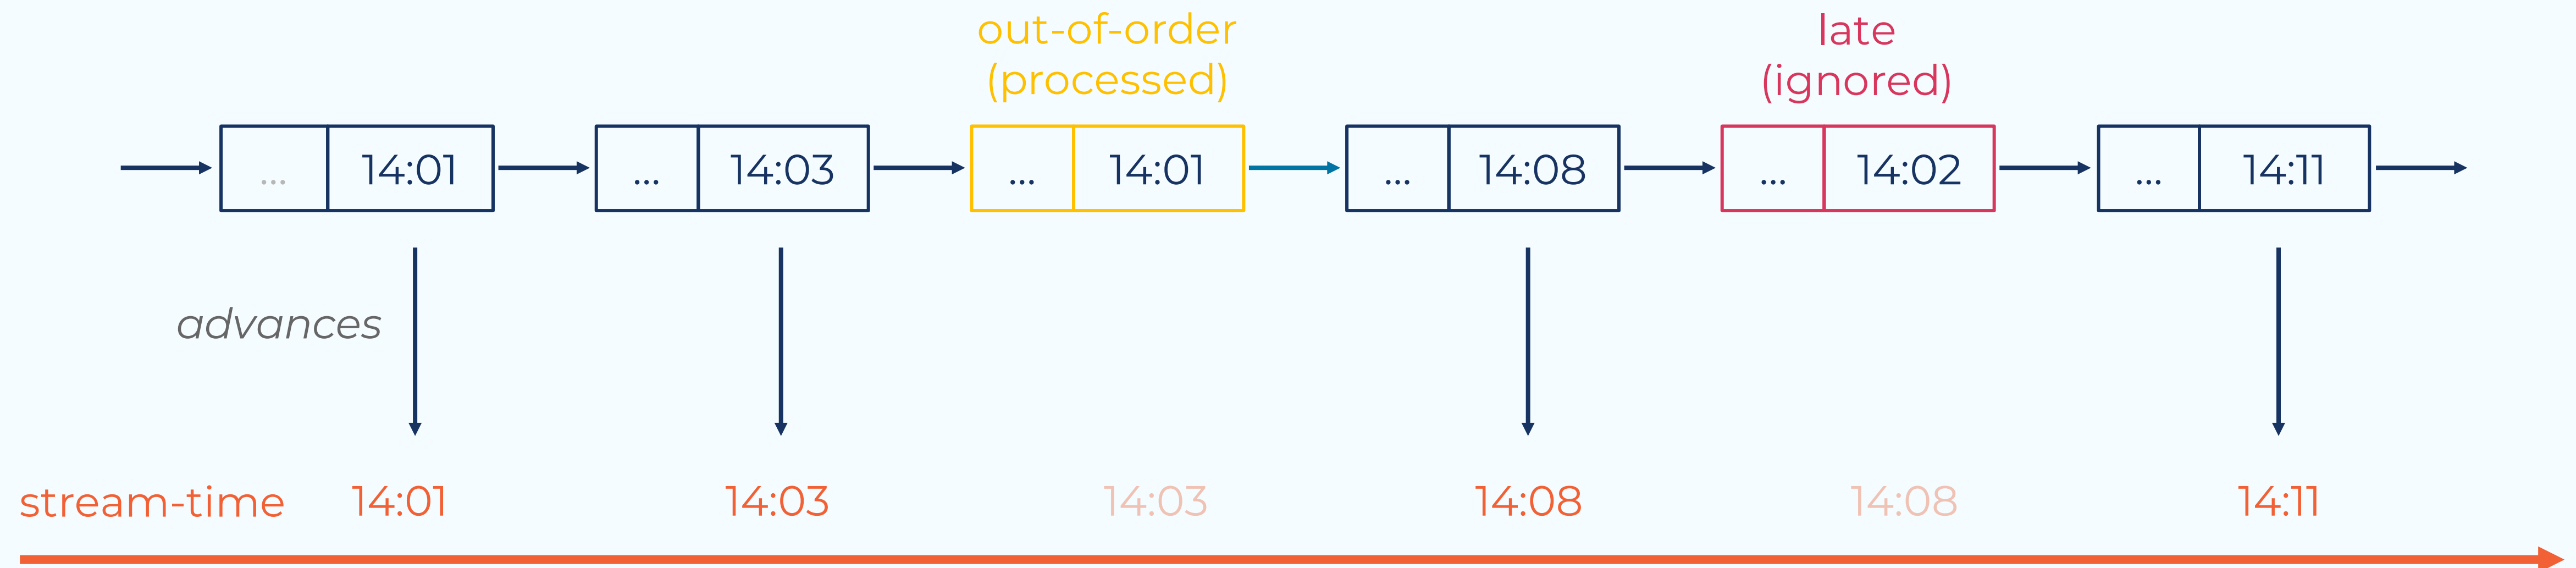
- Any input events with an **event-time** < **stream-time** are considered **out-of-order**
 - For windowed operations, this means the event-timestamp is less than the current stream-time, but is *within* the window time (window size plus grace period)
- Out-of-order records are accepted and processed



[1] cf. [Beyond Analytics: The Evolution of Stream Processing Systems \(SIGMOD 2020\)](#), [Aurora: a new model and architecture for data stream management \(VLDB Journal 2003\)](#)

Time Concepts - Late Input

- The grace period, a *per-window* setting, defines a **cut-off** for **out-of-order events**
- Any out-of-order events that arrive after the grace period are considered **(too) late**, and thus are ignored and not processed
- The delay of an event is determined by **stream-time** - **event-timestamp**



[1] cf. [Beyond Analytics: The Evolution of Stream Processing Systems \(SIGMOD 2020\)](#), [Aurora: a new model and architecture for data stream management \(VLDB Journal 2003\)](#)

Hands On: Time Concepts

The bottom half of the slide features several thick, wavy, light blue lines that sweep upwards from the left towards the right, creating a sense of movement and depth against the dark blue background.

Hands On: Time Concepts

Processor API

The background of the slide is a solid dark blue. In the lower half, there are several thick, wavy, light blue lines that sweep from the left side towards the right, creating a sense of motion and depth.

Processor API

- The DSL provides high-level operators then builds the topology for you
- For any options the DSL doesn't allow you can use Processor API
- Maximum flexibility
 - The developer is responsible for all the details
- Can build any type of stream processing application

Processor API

- Access to state stores for custom stateful operations
- You can call commit from processors
- Provides the ability to schedule arbitrary operations with Punctuation
 - Stream-time processing
 - Wall-clock time

Processor API

- Building streams applications with the Processor API follows this pattern
 - Add source node(s)
 - Add N number of processors that are child nodes of the source node(s)
 - Optionally create one or more StoreBuilder instances and attach them to the processing nodes
 - Add sink node(s) and make them child node of processor nodes

Processor API

- As you add nodes you'll use the name of one node as the parent name for another node
- Processors can have more than one parent node
- The Processor API gives you the flexibility to forward records to all child nodes or just select ones
- We'll see more details in the exercise for this module

Processor API - Example

```
Topology topology = new Topology();
```

```
topology.addSource("source-node", "topicA", "topicB");
```



```
topology.addProcessor("custom-processor", new CustomProcessorSupplier(storeName), "source-node");
```



```
topology.addSink("sink-node", "output-topic", "custom-processor");
```

Processor API

- It's possible to “mix-in” the Processor API with the Streams DSL.
 - transform
 - transformValues
 - process
- This means you can e.g. the more convenient DSL for most processing steps of your application, and only need to use the Processor API for a few more complicated steps.

Hands On: Processor API

The bottom half of the slide features several thick, wavy, light blue lines that sweep upwards from the left towards the right, creating a sense of motion and modern design.

Hands On: Processor API

Testing

The background of the slide features a dark navy blue field. In the lower half, there are several thick, flowing, wavy lines in a lighter blue color. These lines originate from the left side and curve upwards and to the right, creating a sense of movement and depth. The word "Testing" is positioned in the upper left area, above the wavy lines.

Testing Kafka Streams

- Testing is a critical part of software development
- Kafka Streams connects to brokers
- Expensive in unit tests to rely on broker connection
- The **TopologyTestDriver** solves the issue of unit testing a topology
- Note that integration testing with a live broker should still be done

*Testing Kafka Streams Topology***TestDriver**

- Build topology as usual including configurations
 - Use **MockSchemaRegistry** with **TopologyTestDriver**
- Instantiate **TopologyTestDriver** and use a **Topology** and **Properties** as constructor parameters
- Create **TestInputTopic** instances
- Call **TestInputTopic.pipeInput** with **KeyValue** objects
 - Overloaded methods allow for providing timestamps, List of records
- Executing **TestInputTopic.pipeInput** will trigger stream-time punctuation
 - Wall clock punctuations will fire only by calling the **advanceWallClockTime** method

Testing - Testable Application

- The Streams DSL has several operators that accept SAM interfaces
- Can easily use lambda expressions
- Downside - can't easily test the lambda expressions in isolation
- Consider to write the concrete implementations of those interfaces
 - Can write single unit tests against ValueMapper, Reducer etc.

Testing - Integration Tests

- Good to have integration tests against a live broker
- Can see how stateful operations behave in real environment
 - **TopologyTestDriver** doesn't have the caching behavior or commits
 - Doesn't write to real topics
- Best choice for brokers in a test is to use **TestContainers**
 - Easily control broker life-cycle
 - Possible to share a container across multiple test classes
 - Improved testing time

Hands On: Testing

The background of the slide is a solid dark blue. In the lower half, there are several thick, wavy, light blue lines that sweep from the left side towards the right, creating a sense of motion and depth. The lines vary in thickness and curvature, with some starting near the bottom left and others higher up, all generally trending upwards and to the right.

Hands On: Testing

Error Handling

The background of the slide is a solid dark blue. In the lower half, there are several thick, wavy, light blue lines that flow from the left side towards the right, creating a sense of movement and depth.

Error Handling

- Errors are inevitable, especially in distributed systems and applications
- But not all errors are worthy of “stop the world”
 - Network partitions will *usually* resolve quickly
 - Change in partition ownership - very transient
- Need to provide a mechanism for gracefully handling errors
 - Recovery automatically when possible
 - Shut down when truly unrecoverable without intervention

Error Handling

- In Kafka Streams there are 3 broad categories where errors may occur
 - Entry - Consuming records - network or deserialization errors
 - Processing of records - not in expected format
 - Exit - Producing records - network or serialization errors
- For these broad categories there are handlers
 - Provide for users deciding to best deal with recoverable errors
 - In some cases the best approach is to acknowledge and continue
 - Other times more prudent to shut down
- We'll cover the handlers Kafka Streams provides in all 3 areas

Error Handling - Consuming

- For errors consuming into Kafka Streams
 - The **DeserializationExceptionHandler** interface
 - Default configuration is the **LogAndFailExceptionHandler**
 - Other option is to use the **LogAndContinueExceptionHandler**
 - Can provide a custom implementation and provide the classname via the Kafka Streams configuration

Error Handling - Producing

- For errors producing from Kafka Streams
 - The **ProductionExceptionHandler** interface
 - Can respond with continue processing or fail
 - Default configuration is the **DefaultProductionExceptionHandler**
 - The default option always returns fail
 - For any other option you'll need to implement your own implementation
 - Only applies to exceptions not handled by Kafka Streams
 - **RecordTooLargeException**

Error Handling - Client Related

- Kafka Streams uses embedded producer and consumer instances
 - Clients can experience intermittent temporary failures
 - Network partition
 - Lead broker changes
 - Clients have their own configurations these situations
 - Tough to get client configurations correct.
 - Optimizing for resilience means blocking by the clients which has adverse effects on the Kafka Streams application
 - Too loose and there's a risk of application shut downs for transient issues

Error Handling - Processing

- For processing errors
 - Exceptions from user logic bubble up and shut down the application
 - Kafka Streams provides the **StreamsUncaughtExceptionHandler**
 - Works for Exceptions not handled by Kafka Streams
 - The implementation you provide has three options
 - Replace the thread
 - Shutdown the individual streams instance
 - Shutdown all streams instances (with the same application id)

Error Handling - Client Related

- Kafka Streams solution is the **task.timeout.config**
 - When clients experience an issue Kafka Stream starts a timer
 - Kafka Streams attempts to make progress with other tasks
 - The task with the failed operation is retried
 - Procedure continues until success or the timeout is reached

Hands On: Error Handling

The bottom half of the slide features several thick, wavy, light blue lines that sweep across the frame from left to right, creating a sense of motion and modern design.

Hands On: Error Handling

Elasticity, Scaling, Parallelism

The background of the slide is a dark navy blue. In the lower half, there are several thick, wavy, light blue lines that flow from the left side towards the right, creating a sense of movement and depth. The lines vary in thickness and curve, with some starting near the bottom left and others higher up, all trending towards the right edge of the frame.

Elasticity, Scaling, Parallelism

- Kafka Streams internally uses the concept of a task for a work unit
 - Driven by the number of input partitions
 - Kafka Streams takes the highest partition count from the source node and that determines the number of tasks
 - Kafka Streams assigns tasks to StreamThread(s) one is the default
 - Can have as many threads there are tasks
 - Threads beyond this count are idle

Elasticity, Scaling, Parallelism

- Each Task represents a Topic-Partition from the source
- By increasing the thread count - automatically increases Kafka Streams processing power.
- Spinning up new Kafka Streams instances (with the same application ID) provides the same increase in throughput
 - Same ceiling applies more instances than tasks are idle - but can be available for failover

Elasticity, Scaling, Parallelism

- Since Kafka Streams uses **KafkaConsumer** instances internally it automatically inherits powerful dynamic scaling properties
 - Consumer group protocol
 - When a member leaves a rebalance reassigns resources to current active members
 - When a new member joins a rebalance pulls resources from existing members and gives them to the new member

Elasticity, Scaling, Parallelism

- The same protocol applies to Kafka Streams
 - For more processing power spin up as many Kafka Streams instances until all tasks are accounted for
 - In times where the level of traffic is reduced then take down Kafka Stream instances and resources are automatically assigned to active applications
- This behavior is completely dynamic - processing continues after a brief delay for the rebalance to complete

Fault Tolerance

The background of the slide is a dark navy blue. In the lower half, there are several thick, wavy, light blue lines that sweep from the left towards the right, creating a sense of motion and depth. The lines vary in thickness and curvature, with some starting near the bottom left and others higher up, all generally trending towards the right side of the frame.

Fault Tolerance

- Kafka Streams has stateless (e.g., map, filter) as well as stateful operations
- For stateful operations, such as aggregations and joins:
 - State stores are either persistent or in-memory
 - Backed by changelog topics for durability
 - Records written to the changelog as they are persisted
 - Losing a machine with a State Store can happen
 - When starting up Kafka Streams will detect a stateful node
 - If the state doesn't exist fully restored from the changelog topic

Fault Tolerance

- In-memory stores don't retain records across restarts
 - Fully restored to from changelog topic
- Changelog topics use compaction
 - Oldest records by key are deleted
 - Safely leaves most recent records for that key

Fault Tolerance

- With a stateful operation restoring full state can take time
 - Kafka Streams offers stand-by tasks
 - Configure **num.standby.replicas** to number greater than the default setting of zero
 - Kafka Streams will designate another application instance as the “standby”
 - The standby instance keeps a mirrored state store in-sync with the original
 - When the primary goes down, the standby takes over with complete or minor restoration - no downtime restoring

Interactive Queries

The background of the slide is a dark navy blue. In the lower half, there are several thick, wavy, light blue lines that sweep from the left towards the right, creating a sense of motion and depth. The lines vary in thickness and curvature, with some starting near the bottom left and others higher up, all trending towards the right edge of the frame.

Interactive Queries

- Typical pattern with event streaming is reporting
 - Dashboard applications
 - Require the streaming system write results to database
 - UI layer queries the database for live views
- Kafka Streams stateful operations, such as aggregations, represent the present “state” of an event stream
- Interactive queries let you query this state as well as query **KTable** instances

Interactive Queries

- Eligible for interactive queries: KTable and aggregations
- To enable Interactive queries:
 - Name the state store via **Materialized** or using the Stores factory
 - Set the **application.serverconfiguration (host:port)**
- Each application instance has metadata for all instances of the same application (**same application.id**)
- The developer needs to provide the serving layer

Interactive Queries

- Kafka Streams does it without the need for an external database by allowing direct, read-only access to state stores and **KTables**
 - Live - ongoing as it's happening
 - No need to write intermediate results - simplifies architecture
 - A materialized view of the operation in real-time
- ksqlDB: There's also the option to use SQL with ksqlDB to query tables.
 - Take a look at our ksqlDB courses to learn more